

Giuseppe Attanasio

POSTDOCTORAL RESEARCHER · INSTITUTO DE TELECOMUNICAÇÕES

Instituto Superior Técnico, Av. Rovisco Pais, 1, 1049-001, Lisboa

🏠 gattanasio.cc | ✉ giuseppeattanasio6@gmail.com

📱 g8a9 | 🌐 giuseppe-attanasio | 🎓 Scholar | bluesky

“Computers aren’t the thing. They are the thing that gets us to the thing.” - Halt And Catch Fire

Bio

I’m a postdoc in the SARDINE group at Instituto de Telecomunicações, Lisbon, headed by scientific director André Martins. My research interests lie at the intersection of interpretability, fairness, and safety. Previously, I was a MilaNLP member, working with Debora Nozza and Dirk Hovy.

I am a computer engineer who cares about the ethical use of AI on people, communities, and society at large, about which I have given talks and published modest science. Some of the things I’m proud of from my past are co-founding a student team and RiTA, an online community connecting passionate people around Italian resources, and getting decently good at pizza-making. I find cross-institution and cross-discipline collaboration fun and formative. As I trust that deep understanding is rooted in hands-on work, I also enjoy coding for my research projects.

I update my [website](#) more frequently than this document, with recent professional news and projects.

Education & Research

University of Lisbon

POSTDOC AT INSTITUTO DE TELECOMUNICAÇÕES

- PI: André Martins

Lisbon, Portugal

Feb. 2024 - ongoing

Bocconi University

POSTDOC AT DEPARTMENT OF COMPUTING SCIENCES

- PI: Debora Nozza
- Advisor: Dirk Hovy

Milan, Italy

Oct. 2022 - Feb. 2024

Politecnico di Torino

PH.D. AT DEPARTMENT OF CONTROL AND COMPUTER ENGINEERING

- Advisor: Elena Baralis

Turin, Italy

Nov. 2018 - Oct. 2022

Politecnico di Torino

MD IN COMPUTER ENGINEERING, DATA SCIENCE TRACK

- Grade: 110/110 *cum Laude*

Turin, Italy

Oct. 2016 - Oct. 2018

Politecnico di Torino

BD IN COMPUTER ENGINEERING

- Grade: 110/110

Turin, Italy

Oct. 2013 - Oct. 2016

Selected Publications

I firmly believe in open-access, reproducible research. Most of my papers are free to read, and code and data are open-source under a permissive license, whenever possible.

Attanasio, Giuseppe, Savoldi, Beatrice, Fucci, Dennis, and Hovy, Dirk. “Twists, Humps, and Pebbles: Multilingual Speech Recognition Models Exhibit Gender Performance Gaps”. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 21318–21340. URL: <https://aclanthology.org/2024.emnlp-main.1188/>

Zaranis, Emmanouil, **Attanasio, Giuseppe**, Agrawal, Sweta, and Martins, Andre. “Watching the Watchers: Exposing Gender Disparities in Machine Translation Quality Estimation”. In: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar. Vienna, Austria: Association for Computational Linguistics, July 2025, pp. 25261–25284. ISBN: 979-8-89176-251-0. URL: <https://aclanthology.org/2025.acl-long.1228/>

Attanasio, Giuseppe, Plaza del Arco, Flor Miriam, Nozza, Debora, and Lauscher, Anne. “A Tale of Pronouns: Interpretability Informs Gender Bias Mitigation for Fairer Instruction-Tuned Machine Translation”. In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Ed. by Houda Bouamor, Juan Pino, and Kalika Bali. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 3996–4014. URL: <https://aclanthology.org/2023.emnlp-main.243>

Attanasio, Giuseppe, Nozza, Debora, Hovy, Dirk, and Baralis, Elena. “Entropy-based Attention Regularization Frees Unintended Bias Mitigation from Lists”. In: *Findings of the Association for Computational Linguistics: ACL 2022*. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 1105–1119. URL: <https://aclanthology.org/2022.findings-acl.88>

Attanasio, Giuseppe, Pastor, Eliana, Di Bonaventura, Chiara, and Nozza, Debora. “ferret: a Framework for Benchmarking Explainers on Transformers”. In: *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations*. Dubrovnik, Croatia: Association for Computational Linguistics, May 2023, pp. 256–266. URL: <https://aclanthology.org/2023.eacl-demo.29>

Delobelle, Pieter, **Attanasio, Giuseppe**, Nozza, Debora, Blodgett, Su Lin, and Talat, Zeerak. “Metrics for What, Metrics for Whom: Assessing Actionability of Bias Evaluation Metrics in NLP”. in: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 21669–21691. URL: <https://aclanthology.org/2024.emnlp-main.1207/>

Cercas Curry, Amanda, **Attanasio, Giuseppe**, Talat, Zeerak, and Hovy, Dirk. “Classist Tools: Social Class Correlates with Performance in NLP”. in: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Lun-Wei Ku, Andre Martins, and Vivek Srikumar. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 12643–12655. URL: <https://aclanthology.org/2024.acl-long.682/>

Bianchi, Federico, Suzgun, Mirac, **Attanasio, Giuseppe**, Rottger, Paul, Jurafsky, Dan, Hashimoto, Tatsunori, and Zou, James. “Safety-Tuned LLaMAs: Lessons From Improving the Safety of Large Language Models that Follow Instructions”. In: *The Twelfth International Conference on Learning Representations*. 2024. URL: <https://openreview.net/forum?id=gT5hALch9z>

Röttger, Paul, Kirk, Hannah, Vidgen, Bertie, **Attanasio, Giuseppe**, Bianchi, Federico, and Hovy, Dirk. “XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models”. In: *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Ed. by Kevin Duh, Helena Gomez, and Steven Bethard. Mexico City, Mexico: Association for Computational Linguistics, June 2024, pp. 5377–5400. URL: <https://aclanthology.org/2024.naacl-long.301/>

Please, head to my [Google Scholar page](https://scholar.google.com/citations?user=giuseppe.attanasio) or visit <https://gattanasio.cc/uploads/pubs.pdf> for the full list.

Grants & Awards

GRANTS

2024	FCT Advanced Computing , Project: MELT: Multilingual Efficient Learning for Speech-Text Models. Funds: €39,000	FCT
2024	EAMT Sponsorship , Project: A User-Centered Benchmark for Usability Assessment of Speech Translation Systems. Funds: €5,200	EAMT
2023	EAMT Sponsorship , Project: GeFMT: Gender-Fair Language in German Machine Translation. Funds: €5,200	EAMT
2023	ISCRA-C , Project: ItaLLM: Advancing Beyond-English Language Models: A Computational Framework for Italian Language Modeling. Funds: 10,200 GPU hours	CINECA
2023	Google Cloud Credits , Project: Gender Bias in Large Language Models. Funds: €4,800	Google

AWARDS

- | | | |
|------|--|---------|
| 2024 | Social Impact Paper Award , Paper: Twists, Humps, and Pebbles: Multilingual Speech Recognition Models Exhibit Gender Performance Gaps | EMNLP |
| 2023 | Best Student Paper Award , Paper: Contrastive Language–Image Pre-training for the Italian | CLiC-it |

Service Work

Workshop Organization

- (TBD), I will co-organize the next edition of the Gender Bias in NLP workshop.
- 2024, I co-organized the third edition of the Third Workshop on Safety for Conversational AI during LREC COLING 2024.
- 2022, I co-organized the first edition of EvalRS: a Rounded Evaluation of Recommender Systems, a data challenge and workshop held during CIKM 2022.

Editorial

- 2024-ongoing: Area Chair: ACL Rolling Review, LREC/COLING
- 2025-ongoing: Senior Area Chair: LREC/COLING.

Peer Reviewing. I reviewed at least one work submitted to the following venues: ACL Rolling Review (2021-onward); ACM KDD SIGKDD Conference On Knowledge Discovery And Data Mining (2020, 2021); ACM SIGMOD/PODS International Conference on Management of Data (2021); IEEE ICDE: IEEE International Conference on Data Engineering (2020); IEEE ICDM: IEEE International Conference on Data Mining (2021); ACM SAC: ACM/SIGAPP Symposium On Applied Computing (2021); EDBT: International Conference on Extending Database Technology (2020, 2021); DaWaK: International Conference on Big Data Analytics and Knowledge Discovery (2019, 2021).

I reviewed at least one work submitted to the following journals: Expert Systems With Applications, Elsevier; Future Generation Computer Systems, Elsevier; Machine Learning With Applications, Elsevier.

Industry and Academic Collaborations

I have engaged in several fruitful research collaborations with industry partners, leading to impactful outcomes in fairness, inclusivity, multimodality, and recommender systems.

- **Fairness in End-to-End Speech Representation Models** — Studied subgroup biases in neural speech models (e.g., Wav2Vec 2.0, HuBERT) to detect underperformance on underrepresented groups. *Entities: Amazon Alexa AI.*
- **Gender Inclusive Communication in Italian** — Developed methods for detecting and rephrasing non-inclusive language in Italian public and academic texts. *Entities: University of Turin, University of Bologna, University of Bergamo.*
- **Contrastive Language-Image Pre-training for Italian** — Fine-tuned CLIP for Italian, producing the first Italian multimodal CLIP model, finalist at HuggingFace Flax/JAX Community Week. *Entities: Bocconi University, University of Milano-Bicocca, University of Groningen.*
- **EvalRS: Rounded Evaluation of Recommender Systems** — Proposed EvalRS, an open-source framework and challenge addressing fairness and behavioral aspects in RecSys evaluation (CIKM 2022). *Entities: Coveo, Coveo Labs, Stanford University, Microsoft, NVIDIA.*
- **Vision and Language for Fashion Concept Learning** — Demonstrated that CLIP-based multi-modal models efficiently capture general fashion concepts (e.g., sleeve length, heel type). *Entities: Coveo, Coveo Labs, Stanford University, Farfetch.*

Teaching

Teaching Assistant. I have been a teaching assistant for several courses at Politecnico di Torino. In total, I have held 171 hours of complementary teaching. The list of courses where I TAd is:

(2021) **Data science lab: processes and methods**, 60h, MD in Data Science and Engineering; (2020) **Data science lab: processes and methods**, 39h, MD in Data Science and Engineering; (2020) **Business Intelligence for Big Data**, 21h, MD in Industrial engineering and management; (2019) ***Data science lab: processes and methods**, 39h, MD in Data Science and Engineering; (2019) **Business Intelligence for Big Data**, 21h, MD in Industrial engineering and management; (2018) **Basi di dati**, 21h, BD in Computer Engineering; (2018) **The Fourth Industrial Revolution: Promises and Pitfalls in Blending New and Traditional Approaches in Manufacturing and Service Sectors**, 30h, Alta Scuola Politecnica School.

* I played a key role in launching Data Science Lab: Process and Methods for the Data Science and Engineering Master's at Politecnico. The course introduces Python and core Data Science and Machine Learning libraries, with 10 labs providing 60+ pages of exercises and 250+ pages of solutions.

Tutoring. I co-tutored the second edition of the **Mediterranean Machine Learning (M2L) Summer School**, with ownership of the Natural Language Processing part. We prepared a shared notebook to guide students through 1) implementing a Transformer from scratch, 2) training the model for language modeling and machine translation, and 3) testing on gender bias. I have also tutored the hands-on laboratories during the 2024 **LxMLS Machine Learning School** in Lisbon. I co-founded **MAchine Learning At poliTO (MALTO)**, a university-funded student team at Politecnico di Torino. The team's primary goal is to participate in various data science competitions. I have supervised the work of more than **10 master's students**.

Guest Lecturer.

- 2023, Bocconi University: **Contrastive Learning**
- 2023, Bocconi University: **Reinforcement Learning from Human Feedback**
- 2023, Tilburg University: **Social Biases in LMs: Why They Are There, How Do We Measure and Mitigate Them**

INDUSTRY & ASSOCIATIONS

I have held the following courses:

- (2024) **Generative Artificial Intelligence and Basic Prompting: How to Train AI Tools for Work in Publishing Houses**
AIE: Associazione Italiana Editori (online)
- (2021) **Data Theory: Data Visualization with Python**
Reply: Digital Services, Technology and Consulting, Turin (online)
- (2020) **Python technologies for Data Analytics**
Applied Mechatronic Engineering & Technologies, Turin

Invited Talks

I have disseminated my research through several invited talks. Some of the host institutions and entities were: University of Padova, the Italian Association of Editors, Cambridge Forum of Science and Humanities, Karger Publishers, Papers We Love (Milan's Chapter), and SmartData@PoliTO.

Work Experience

Kupata S.r.l.

FOUNDER

Turin, Italy

Nov. 2016 - 2020

- Kupata's main goal is to streamline the Lost and Found process. It brings innovation with a solution that helps people in returning items in a simple, secure, and undisclosed way. The business involves a physical object, the Kupa, and a social community that encourages members to act in the right way.

Consoft Sistemi S.p.A.

CURRICULAR INTERN

Turin, Italy

Mar. 2016 - Jul. 2016

- I participated to the bootstrap phase of the Consoft's proprietary Knowledge Base platform. We built our solution upon Orange HRM, an open source PHP-based platform.