

The following is the complete list of publications (last update: September 13, 2025).

Please, head to my [Google Scholar page](#) for citation counts.

Savoldi, Beatrice, **Attanasio, Giuseppe**, Cupin, Eleonora, Gkovedarou, Eleni, Hackenbuchner, Janiça, Lauscher, Anne, Negri, Matteo, Piergentili, Andrea, Thind, Manjinder, and Bentivogli, Luisa. “Mind the Inclusivity Gap: Multilingual Gender-Neutral Translation Evaluation with mGeNTE”. in: *The 2025 Conference on Empirical Methods in Natural Language Processing (proceedings not out yet)*. 2025. URL: <https://openreview.net/forum?id=dBUHC2QyBh>

Pranav, A, Hackenbuchner, Janiça, **Attanasio, Giuseppe**, Lardelli, Manuel, and Lauscher, Anne. “Glitter: A Multi-Sentence, Multi-Reference Benchmark for Gender-Fair German Machine Translation”. In: *The 2025 Conference on Empirical Methods in Natural Language Processing (proceedings not out yet)*. 2025. URL: <https://openreview.net/forum?id=k9GgDddCUD>

Zaranis, Emmanouil, **Attanasio, Giuseppe**, Agrawal, Sweta, and Martins, Andre. “Watching the Watchers: Exposing Gender Disparities in Machine Translation Quality Estimation”. In: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar. Vienna, Austria: Association for Computational Linguistics, July 2025, pp. 25261–25284. ISBN: 979-8-89176-251-0. DOI: 10.18653/v1/2025.acl-long.1228. URL: <https://aclanthology.org/2025.acl-long.1228/>

Mitchell, Margaret, **Attanasio, Giuseppe**, Baldini, Ioana, Clinciu, Miruna, Clive, Jordan, De-labelle, Pieter, Dey, Manan, Hamilton, Sil, Dill, Timm, Doughman, Jad, Dutt, Ritam, Ghosh, Avijit, Forde, Jessica Zosa, Holtermann, Carolin, Kaffee, Lucie-Aimée, Laud, Tanmay, Lauscher, Anne, Lopez-Davila, Roberto L, Masoud, Maraim, Nangia, Nikita, Ovalle, Anaelia, Pistilli, Gida, Radev, Dragomir, Savoldi, Beatrice, Raheja, Vipul, Qin, Jeremy, Ploeger, Esther, Subramonian, Arjun, Dhole, Kaustubh, Sun, Kaiser, Djanibekov, Amirbek, Mansurov, Jonibek, Yin, Kayo, Cueva, Emilio Villa, Mukherjee, Sagnik, Huang, Jerry, Shen, Xudong, Gala, Jay, Al-Ali, Hamdan, Djanibekov, Tair, Mukhituly, Nurdaulet, Nie, Shangrui, Sharma, Shanya, Stanczak, Karolina, Szczechla, Eliza, Timponi Torrent, Tiago, Tunuguntla, Deepak, Viridiano, Marcelo, Van Der Wal, Oskar, Yakefu, Adina, Névél, Aurélie, Zhang, Mike, Zink, Sydney, and Talat, Zeerak. “SHADES: Towards a Multilingual Assessment of Stereotypes in Large Language Models”. In: *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Ed. by Luis Chiruzzo, Alan Ritter, and Lu Wang. Albuquerque, New Mexico: Association for Computational Linguistics, Apr. 2025, pp. 11995–12041. ISBN: 979-8-89176-189-6. DOI: 10.18653/v1/2025.naacl-long.600. URL: <https://aclanthology.org/2025.naacl-long.600/>

Fucci, Dennis, Gaido, Marco, Negri, Matteo, Bentivogli, Luisa, Martins, Andre, and **Attanasio, Giuseppe**. “Different Speech Translation Models Encode and Translate Speaker Gender Differently”. In: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. Ed. by Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar. Vienna, Austria: Association for Computational Linguistics, July 2025, pp. 1005–1019. ISBN: 979-8-89176-252-7. DOI: 10.18653/v1/2025.acl-short.78. URL: <https://aclanthology.org/2025.acl-short.78/>

Lardelli, Manuel, **Attanasio, Giuseppe**, and Lauscher, Anne. “Building Bridges: A Dataset for Evaluating Gender-Fair Machine Translation into German”. In: *Findings of the Association for Computational Linguistics: ACL 2024*. Ed. by Lun-Wei Ku, Andre Martins, and Vivek Srikumar. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 7542–7550. DOI: 10.18653/v1/2024.findings-acl.448. URL: <https://aclanthology.org/>

2024.findings-acl.448/

Delobelle, Pieter, **Attanasio, Giuseppe**, Nozza, Debora, Blodgett, Su Lin, and Talat, Zeerak. “Metrics for What, Metrics for Whom: Assessing Actionability of Bias Evaluation Metrics in NLP”. in: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 21669–21691. DOI: 10.18653/v1/2024.emnlp-main.1207. URL: <https://aclanthology.org/2024.emnlp-main.1207/>

Röttger, Paul, Kirk, Hannah, Vidgen, Bertie, **Attanasio, Giuseppe**, Bianchi, Federico, and Hovy, Dirk. “XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models”. In: *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Ed. by Kevin Duh, Helena Gomez, and Steven Bethard. Mexico City, Mexico: Association for Computational Linguistics, June 2024, pp. 5377–5400. DOI: 10.18653/v1/2024.naacl-long.301. URL: <https://aclanthology.org/2024.naacl-long.301/>

Attanasio, Giuseppe, Savoldi, Beatrice, Fucci, Dennis, and Hovy, Dirk. “Twists, Humps, and Pebbles: Multilingual Speech Recognition Models Exhibit Gender Performance Gaps”. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 21318–21340. DOI: 10.18653/v1/2024.emnlp-main.1188. URL: <https://aclanthology.org/2024.emnlp-main.1188/>

Cercas Curry, Amanda, **Attanasio, Giuseppe**, Talat, Zeerak, and Hovy, Dirk. “Classist Tools: Social Class Correlates with Performance in NLP”. in: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Lun-Wei Ku, Andre Martins, and Vivek Srikumar. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 12643–12655. DOI: 10.18653/v1/2024.acl-long.682. URL: <https://aclanthology.org/2024.acl-long.682/>

Pastor, Eliana, Koudounas, Alkis, **Attanasio, Giuseppe**, Hovy, Dirk, and Baralis, Elena. “Explaining Speech Classification Models via Word-Level Audio Segments and Paralinguistic Features”. In: *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Yvette Graham and Matthew Purver. St. Julian’s, Malta: Association for Computational Linguistics, Mar. 2024, pp. 2221–2238. DOI: 10.18653/v1/2024.eacl-long.136. URL: <https://aclanthology.org/2024.eacl-long.136/>

Koudounas, Alkis, Pastor, Eliana, **Attanasio, Giuseppe**, Mazzia, Vittorio, Giollo, Manuel, Guedre, Thomas, Reale, Elisa, Cagliero, Luca, Cumani, Sandro, Alfaro, Luca de, et al. “Towards comprehensive subgroup performance analysis in speech models”. In: *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 32 (2024), pp. 1468–1480

Koudounas, Alkis, Pastor, Eliana, **Attanasio, Giuseppe**, Alfaro, Luca de, and Baralis, Elena. “Prioritizing data acquisition for end-to-end speech model improvement”. In: *ICASSP 2024-2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE. 2024, pp. 7000–7004

Attanasio, Giuseppe, Plaza del Arco, Flor Miriam, Nozza, Debora, and Lauscher, Anne. “A Tale of Pronouns: Interpretability Informs Gender Bias Mitigation for Fairer Instruction-Tuned Machine Translation”. In: *Proceedings of the 2023 Conference on Empirical Methods in Natural*

Language Processing. Ed. by Houda Bouamor, Juan Pino, and Kalika Bali. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 3996–4014. DOI: 10.18653/v1/2023.emnlp-main.243. URL: <https://aclanthology.org/2023.emnlp-main.243>

Attanasio, Giuseppe, Pastor, Eliana, Di Bonaventura, Chiara, and Nozza, Debora. “ferret: a Framework for Benchmarking Explainers on Transformers”. In: *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations*. Dubrovnik, Croatia: Association for Computational Linguistics, May 2023, pp. 256–266. URL: <https://aclanthology.org/2023.eacl-demo.29>

Cassani, Giovanni, Bianchi, Federico, **Attanasio, Giuseppe**, Marelli, Marco, and Guenther, Fritz. *Meaning Modulations and Stability in Large Language Models: An Analysis of BERT Embeddings for Psycholinguistic Research*. Oct. 2023. DOI: 10.31234/osf.io/b45ys. URL: osf.io/preprints/psyarxiv/b45ys_v1

Chia, Patrick John, **Attanasio, Giuseppe**, Tagliabue, Jacopo, Bianchi, Federico, Greco, Ciro, Moreira, Gabriel de Souza P., Eynard, Davide, and Husain, Fahd. *E Pluribus Unum: Guidelines on Multi-Objective Evaluation of Recommender Systems*. arXiv:2304.10621 [cs]. Apr. 2023. DOI: 10.48550/arXiv.2304.10621. URL: <http://arxiv.org/abs/2304.10621> (visited on 09/15/2023)

Cercas Curry, Amanda, **Attanasio, Giuseppe**, Nozza, Debora, and Hovy, Dirk. “MilaNLP at SemEval-2023 Task 10: Ensembling Domain-Adapted and Regularized Pretrained Language Models for Robust Sexism Detection”. In: *Proceedings of the 17th International Workshop on Semantic Evaluation (SemEval-2023)*. Toronto, Canada: Association for Computational Linguistics, July 2023, pp. 2067–2074. DOI: 10.18653/v1/2023.semeval-1.285. URL: <https://aclanthology.org/2023.semeval-1.285> (visited on 09/15/2023)

Bonaldi, Helena, **Attanasio, Giuseppe**, Nozza, Debora, and Guerini, Marco. “Weigh Your Own Words: Improving Hate Speech Counter Narrative Generation via Attention Regularization”. In: *Proceedings of the 1st Workshop on CounterSpeech for Online Abuse (CS4OA)*. ed. by Yi-Ling Chung, Helena Bonaldi, Gavin Abercrombie, and Marco Guerini. Prague, Czechia: Association for Computational Linguistics, Sept. 2023, pp. 13–28. URL: <https://aclanthology.org/2023.cs4oa-1.2/>

Koudounas, Alkis, Pastor, Eliana, **Attanasio, Giuseppe**, Mazzia, Vittorio, Giollo, Manuel, Guedre, Thomas, Cagliero, Luca, Alfaro, Luca de, Baralis, Elena, and Amberti, Daniele. “Exploring subgroup performance in End-to-End speech models”. In: *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE. 2023

Koudounas, Alkis, La Quatra, Moreno, Vaiani, Lorenzo, Colomba, Luca, **Attanasio, Giuseppe**, Pastor, Eliana, Cagliero, Luca, and Baralis, Elena. “ITALIC: An ITALian Intent Classification Dataset”. In: *INTERSPEECH 2023*. IEEE. 2023

Chia, Patrick John, **Attanasio, Giuseppe**, Bianchi, Federico, Terragni, Silvia, Magalhães, Ana Rita, Goncalves, Diogo, Greco, Ciro, and Tagliabue, Jacopo. “Fashionclip: Connecting language and images for product representations”. In: *arXiv preprint arXiv:2204.03972* (2022)

Pipoli, Vittorio, **Attanasio, Giuseppe**, Lovino, Marta, and Ficarra, Elisa. “Squeeze and Learn: Compressing Long Sequences with Fourier Transformers for Gene Expression Prediction”. In: (2023)

Attanasio, Giuseppe, Nozza, Debora, and Bianchi, Federico. “MilaNLP at SemEval-2022 Task 5: Using Perceiver IO for Detecting Misogynous Memes with Text and Image Modalities”. In: *Proceedings of the 16th International Workshop on Semantic Evaluation (SemEval-2022)*. Seattle,

United States: Association for Computational Linguistics, July 2022, pp. 654–662. DOI: 10.18653/v1/2022.semeval-1.90. URL: <https://aclanthology.org/2022.semeval-1.90>

Nozza, Debora, Bianchi, Federico, and **Attanasio, Giuseppe**. “HATE-ITA: Hate Speech Detection in Italian Social Media Text”. In: *Proceedings of the Sixth Workshop on Online Abuse and Harms (WOAH)*. Seattle, Washington (Hybrid): Association for Computational Linguistics, July 2022, pp. 252–260. DOI: 10.18653/v1/2022.woah-1.24. URL: <https://aclanthology.org/2022.woah-1.24>

Attanasio, Giuseppe, Nozza, Debora, Pastor, Eliana, and Hovy, Dirk. “Benchmarking Post-Hoc Interpretability Approaches for Transformer-based Misogyny Detection”. In: *Proceedings of NLP Power! The First Workshop on Efficient Benchmarking in NLP*. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 100–112. DOI: 10.18653/v1/2022.nlppower-1.11. URL: <https://aclanthology.org/2022.nlppower-1.11>

Attanasio, Giuseppe, Nozza, Debora, Hovy, Dirk, and Baralis, Elena. “Entropy-based Attention Regularization Frees Unintended Bias Mitigation from Lists”. In: *Findings of the Association for Computational Linguistics: ACL 2022*. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 1105–1119. DOI: 10.18653/v1/2022.findings-acl.88. URL: <https://aclanthology.org/2022.findings-acl.88>

Bellocca, Gian Pietro, **Attanasio, Giuseppe**, Cagliero, Luca, and Fior, Jacopo. “Leveraging the momentum effect in machine learning-based cryptocurrency trading”. In: *Machine Learning with Applications* 8 (2022), p. 100310

Tagliabue, Jacopo, Bianchi, Federico, Schnabel, Tobias, **Attanasio, Giuseppe**, Greco, Ciro, Moreira, Gabriel de Souza P, and Chia, Patrick John. “EvalRS: a Rounded Evaluation of Recommender Systems”. In: *arXiv preprint arXiv:2207.05772* (2022)

Attanasio, Giuseppe, Nozza, Debora, Bianchi, Federico, and Hovy, Dirk. “Is It Worth the (Environmental) Cost? Limited Evidence for the Benefits of Diachronic Continuous Training”. In: *arXiv preprint arXiv:2210.07365* (2022)

Cassani, Giovanni, **Attanasio, Giuseppe**, and Bianchi, Federico. “Sound symbolic associations via contrastive learning of text and images”. English. In: *Architectures and Mechanisms for Language Processing* (2022), AMLaP ; Conference date: 07-09-2022 Through 09-09-2022. Sept. 2022, pp. 1–2. URL: <https://amlap2022.york.ac.uk>

Chia, Patrick John, **Attanasio, Giuseppe**, Bianchi, Federico, Terragni, Silvia, Magalhães, Ana Rita, Goncalves, Diogo, Greco, Ciro, and Tagliabue, Jacopo. “Contrastive Language and Vision Learning of General Fashion Concepts”. In: *Scientific reports* (2022)

Bianchi, Federico, **Attanasio, Giuseppe**, Pisoni, Raphael, Terragni, Silvia, Sarti, Gabriele, and Lakshmi, Sri. “Contrastive Language-Image Pre-training for the Italian Language”. In: *arXiv:2108.08688 [cs]* (Aug. 2021). arXiv: 2108.08688. URL: <http://arxiv.org/abs/2108.08688> (visited on 08/21/2021)

Attanasio, Giuseppe, Greco, Salvatore, La Quatra, Moreno, Cagliero, Luca, Tonti, Michela, Cerquitelli, Tania, and Raus, Rachele. “E-MIMIC: Empowering Multilingual Inclusive Communication”. In: *Proceedings of the First International Workshop on Data science for equality, inclusion and well-being challenges*. Dec. 2021. URL: <https://iris.polito.it/handle/11583/2946252> (visited on 12/21/2021)

Cagliero, Luca, Garza, Paolo, **Attanasio, Giuseppe**, and Baralis, Elena. “Training ensembles of faceted classification models for quantitative stock trading”. en. In: *Computing* 102.5 (May

2020), pp. 1213–1225. ISSN: 1436-5057. DOI: 10.1007/s00607-019-00776-7. URL: <https://doi.org/10.1007/s00607-019-00776-7> (visited on 11/15/2020)

Attanasio, G., Giobergia, F., Pasini, A., Ventura, F., Baralis, E., Cagliero, L., Garza, P., Apiletti, D., Cerquitelli, T., and Chiusano, S. “DSLE: A Smart Platform for Designing Data Science Competitions”. In: *2020 IEEE 44th Annual Computers, Software, and Applications Conference (COMPSAC)*. ISSN: 0730-3157. July 2020, pp. 133–142. DOI: 10.1109/COMPSAC48688.2020.00026

Attanasio, Giuseppe, Cagliero, Luca, and Baralis, Elena. “Leveraging the explainability of associative classifiers to support quantitative stock trading”. In: *Proceedings of the Sixth International Workshop on Data Science for Macro-Modeling*. DSMM ’20. New York, NY, USA: Association for Computing Machinery, June 2020, pp. 1–6. ISBN: 978-1-4503-8030-0. DOI: 10.1145/3401832.3402679. URL: <https://doi.org/10.1145/3401832.3402679> (visited on 11/15/2020)

Attanasio, Giuseppe and Pastor, Eliana. “PoliTeam @ AMI: Improving Sentence Embedding Similarity with Misogyny Lexicons for Automatic Misogyny Identification in Italian Tweets”. en. In: *EVALITA Evaluation of NLP and Speech Tools for Italian - December 17th, 2020*. Ed. by Valerio Basile, Danilo Croce, Maria Maro, and Lucia C. Passaro. Accademia University Press, 2020, pp. 48–54. ISBN: 9791280136329. DOI: 10.4000/books.aaccademia.6807. URL: <http://books.openedition.org/aaccademia/6807> (visited on 08/16/2021)

Attanasio, Giuseppe, Cagliero, Luca, Garza, Paolo, and Baralis, Elena. “Quantitative cryptocurrency trading: exploring the use of machine learning techniques”. In: *Proceedings of the 5th Workshop on Data Science for Macro-modeling with Financial and Economic Datasets*. DSMM’19. New York, NY, USA: Association for Computing Machinery, June 2019, pp. 1–6. ISBN: 978-1-4503-6823-0. DOI: 10.1145/3336499.3338003. URL: <https://doi.org/10.1145/3336499.3338003> (visited on 11/15/2020)

Attanasio, G., Cagliero, L., Garza, P., and Baralis, E. “Combining News Sentiment and Technical Analysis to Predict Stock Trend Reversal”. In: *2019 International Conference on Data Mining Workshops (ICDMW)*. ISSN: 2375-9259. Nov. 2019, pp. 514–521. DOI: 10.1109/ICDMW.2019.00079

Attanasio, G., Cannavò, A., Cibrario, F., Lamberti, F., Montuschi, P., and Paravati, G. “HOT: Hold your own tools for AR-based constructive art”. In: *2017 IEEE Symposium on 3D User Interfaces (3DUI)*. Mar. 2017, pp. 256–257. DOI: 10.1109/3DUI.2017.7893369